


Chapter 1: Introduction (1)





Introduction

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Operations
- Resource Management
- Security and Protection
- Virtualization
- Distributed Systems
- Kernel Data Structures
- Computing Environments
- Free/Libre and Open-Source Operating Systems





Objectives

- Describe the general organization of a computer system and the role of interrupts
- Describe the components in a modern, multiprocessor computer system
- Illustrate the transition from user mode to kernel mode
- Discuss how operating systems are used in various computing environments
- Provide examples of free and open-source operating systems

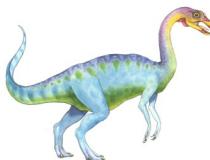




What Does Operating System Mean

- An operating system is “fill in the blanks”
- What about:
 - Car
 - Airplane
 - Printer
 - Washing Machine
 - Toaster
 - Compiler
 -





What is an Operating System


- A program that acts as an intermediary between a user of a computer and the computer hardware, including the following functions
 - Process management
 - Memory management
 - File management
 - I/O management
- Operating system goals
 - Execute user programs and make the solving of user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner





Computer System Structure

- Computer system can be divided into four components:
 - Hardware – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - Operating system
 - ▶ Controls and coordinates use of hardware among various applications and users
 - Application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games, etc.
 - Users - People, machines, other computers





What Operating Systems Do

- Depends on the point of view
 - Users want convenience, **ease of use** and **good performance**
 - ▶ Don't care about **resource utilization**
 - The shared computer such as **mainframe** or **minicomputer** must keep all users happy
 - ▶ OS is a **resource allocator** and **control program** making efficient use of HW and managing execution of user programs
 - Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
 - Mobile devices like smartphones and tablets are resource poor, optimized for usability and battery life
 - ▶ Mobile user interfaces such as touch screens, voice recognition
 - Some computers have little or no user interface, such as embedded computers in devices and automobiles
 - ▶ Run primarily without user intervention





Operating System Definition


- No universally accepted definition
 - Everything a vendor ships when you order an operating system is a good approximation, but varies wildly
- The one program running at all times on the computer is the **kernel**, part of the operating system
- Everything else is either
 - A **system program** (ships with the operating system, but not part of the kernel) , or
 - An **application program**, all programs not associated with the operating system
- Today's OSes for general purpose and mobile computing also include **middleware** – a set of software frameworks that provide additional services to application developers such as databases, multimedia, graphics






Computer System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common **bus** providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles





Computer-System Operations

- I/O devices and the CPU can execute concurrently
 - Each device controller is in charge of a particular device type
 - Each device controller has a local buffer
 - Each device controller type has an operating system **device driver** to manage it
 - CPU moves data from main memory to local buffers of device controllers and vice versa
 - I/O is from the device to local buffer of controller
 - Device controller informs CPU that it has finished its operation by causing an **interrupt**






Interrupt Handling

- The operating system preserves the state of the CPU by storing the registers and the program counter
- Determines which type of interrupt has occurred
- Separate segments of code determine what action should be taken for each type of interrupt
- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- An operating system is **interrupt driven**





Interrupt-Drive I/O Cycle





I/O Structure

- Two methods for I/O handling
 - Synchronous - After I/O starts, control returns to user program only upon I/O completion
 - ▶ Wait instruction idles the CPU until the next interrupt
 - ▶ Wait loop (contention for memory access)
 - ▶ At most one I/O request is outstanding at a time, no simultaneous I/O processing
 - Asynchronous - After I/O starts, control returns to user program without waiting for I/O completion
 - ▶ **System call** – request to the OS to allow user to wait for I/O completion
 - ▶ **Device-status table** contains entry for each I/O device indicating its type, address, and state
 - ▶ OS indexes into I/O device table to determine device status and to modify table entry to include interrupt





Computer Startup

- **Bootstrap loader** - A small program that places the operating system of a computer into memory
 - Typically stored in ROM or EPROM, generally known as **firmware**
 - When a computer powers up or restarts, the BIOS performs some initial tests and then transfers control to the master boot record, where the bootstrap loader resides
 - Initialize all aspects of a system
 - Load operating system kernel and starts execution
 - Most new computers ship with boot loaders for some version of Windows or macOS

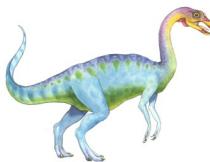




Storage Structure (1)


- Main memory – only large storage media that the CPU can access directly
 - **Random access**
 - **Volatile**
 - In the form of **Dynamic Random-Access Memory (DRAM)**
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity





Storage Structure (2)

- **Hard Disk Drives (HDD)** – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer
- **Non-volatile memory (NVM)** devices– faster than hard disks, nonvolatile
 - Various technologies
 - Become more popular as capacity and performance increases, price drops






Storage Definitions and Notation Review

The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes. A **kilobyte**, or **KB**, is 1,024 bytes; a **megabyte**, or **MB**, is $1,024^2$ bytes; a **gigabyte**, or **GB**, is $1,024^3$ bytes; a **terabyte**, or **TB**, is $1,024^4$ bytes; and a **petabyte**, or **PB**, is $1,024^5$ bytes. Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).






Storage Hierarchy

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage
- **Device Driver** for each device controller to manage I/O
 - Provides uniform interface between controller and kernel






Storage-Device Hierarchy



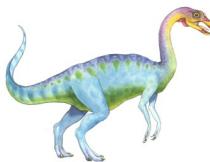


How a Modern Computer Works




A Von Neumann architecture





Direct Memory Access Structure

- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte





Operating System Operations

- Bootstrap program
 - A simple code to initialize the system, and load the kernel
- Starts **system daemons** (services provided outside of the kernel)
- Kernel **interrupt driven** (hardware and software)
 - Hardware interrupt by one of the devices
 - Software interrupt (**exception** or **trap**)
 - ▶ Software error (e.g., division by zero)
 - ▶ Request for operating system service – **system call**
 - ▶ Other process problems include infinite loop, processes modifying each other or the operating system





Multiprogramming (Batch System)

- Single user cannot always keep CPU and I/O devices busy
- Multiprogramming organizes jobs (code and data) so CPU always has one to execute
- A subset of total jobs in system is kept in memory
- One job selected and run via **job scheduling**
- When job has to wait (for I/O for example), OS switches to another job






Multitasking (Timesharing)

- A logical extension of Batch systems— the CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - **Response time** should be < 1 second
 - Each user has at least one program executing in memory ⇒ **process**
 - If several jobs ready to run at the same time ⇒ **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out to run
 - **Virtual memory** allows execution of processes not completely in memory





Memory Layout for Multiprogrammed System





Dual-Mode Operation (1)


- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
- **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code.
 - When a user is running \Rightarrow mode bit is “user”
 - When kernel code is executing \Rightarrow mode bit is “kernel”





Dual-Mode Operation (2)

- How do we guarantee that user does not explicitly set the mode bit to “kernel”?
 - System call changes mode to kernel, return from call resets it to user
- Some instructions designated as **privileged**, only executable in kernel mode





Timer

- Timer to prevent infinite loop (or process hogging resources)
 - Timer is set to interrupt the computer after some time period
 - Keep a counter that is decremented by the physical clock
 - Operating system set the counter (privileged instruction)
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time



Take a Break

