



Operating System (CSC 3150)

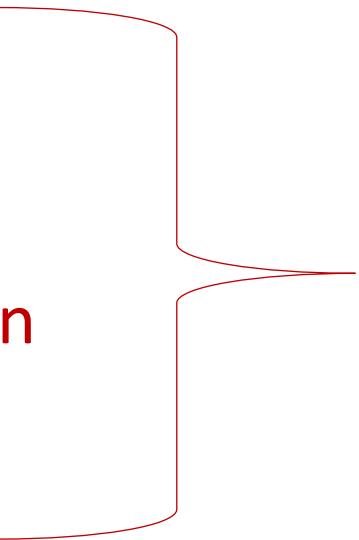
Tutorial 4

Liangsen Wang:
224040364@link.cuhk.edu.cn

Target

In this tutorial, we will practice **Pthread** programming using c/c++.

- Process
- Thread
- Pthread creation
- Pthread termination
- Pthread join
- Pthread mutex
- Pthread conditional value



Assignment 2

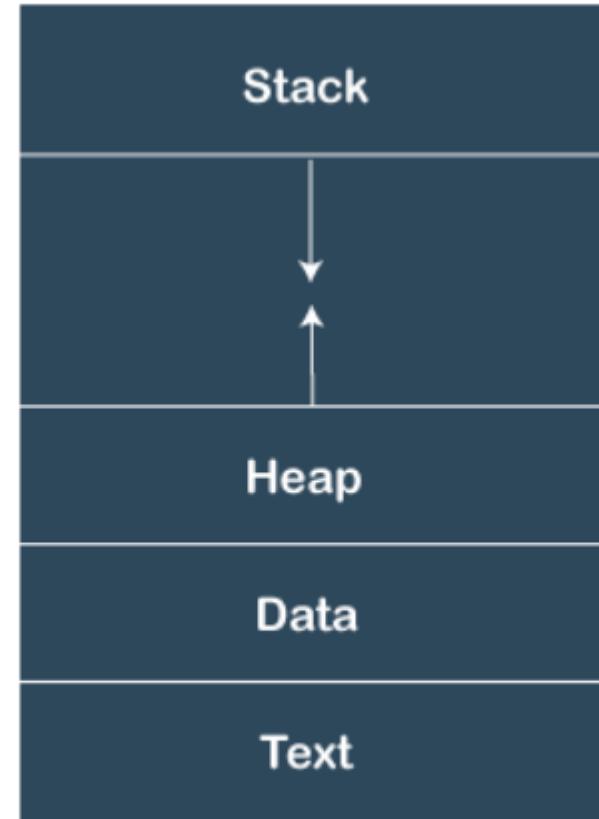
What is Process?

- A process is **an instance of a program** that is being executed.

Features of Process

- Each time we create a process, we need to make a separate system call for each process to the OS. The **fork()** function creates the process.
- **Each process exists within its own address or memory space.**
- Each process **is independent and treated as an isolated process by the OS.**
- Processes need IPC (Inter-process Communication) in order to communicate with each other.

max

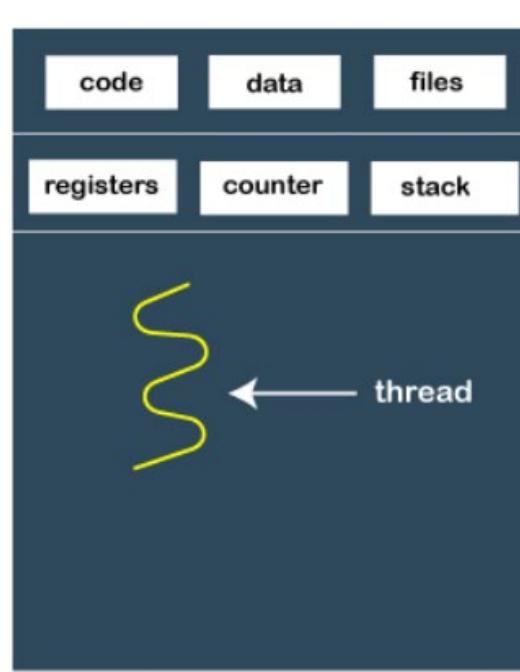


Process in the Memory

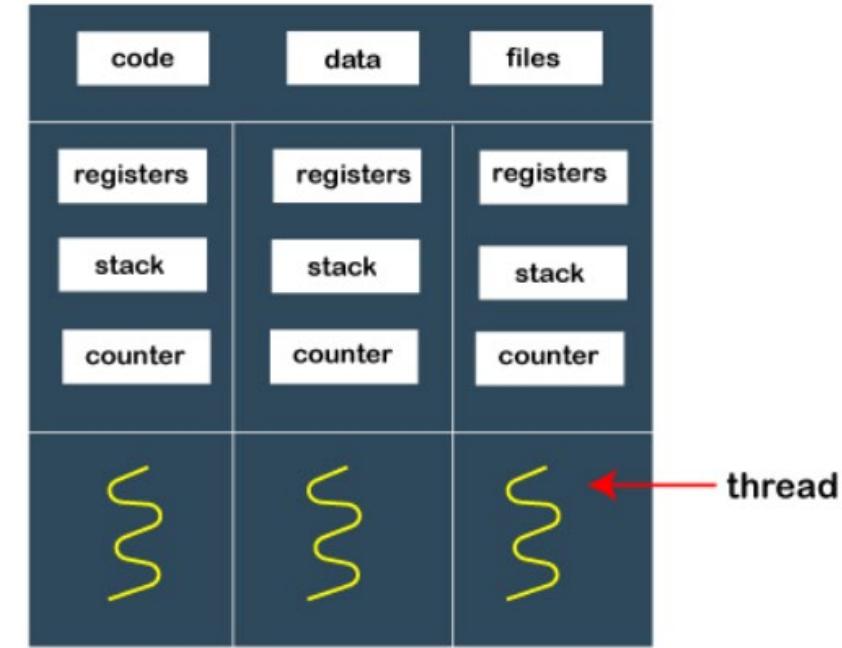
What is a thread?

A **thread** is a single sequence stream within a process. Because threads have some of the properties of processes, they are sometimes called *lightweight processes*.

- Threads are not independent from each other unlike processes. As a result, **threads shares with other threads their code section, data section and OS resources like open files and signals**. But, like processes, a thread has its own **program counter (PC)**, a **register set**, and a **stack space**.
- Threads use and exist within these process resources. They are able to be scheduled by the operating system and run as independent entities within a process.
- A process can have multiple threads**, all of which share the resources within a process and all of which execute within the same address space.



Single-threaded process



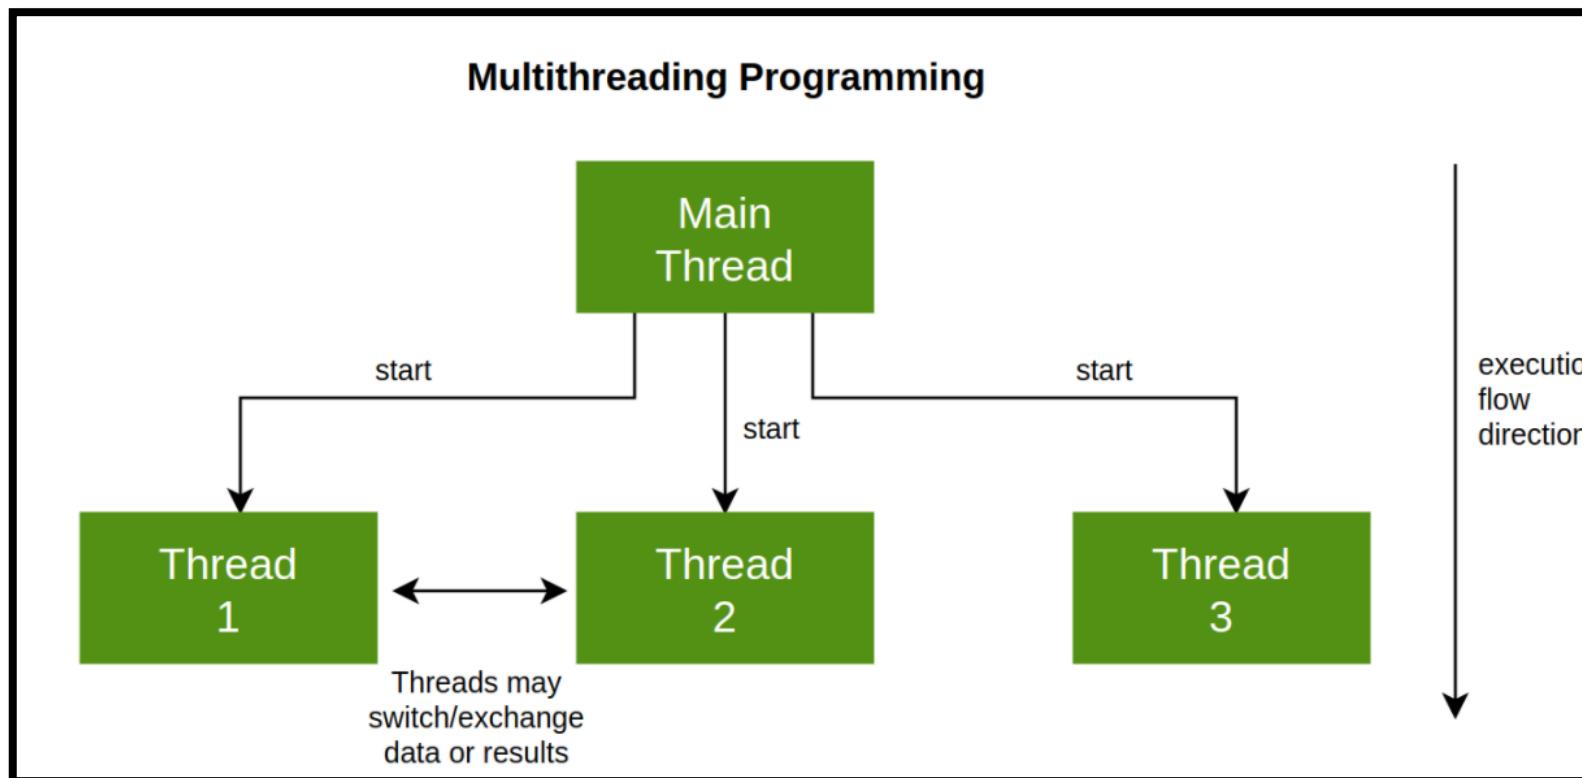
Multi-threaded process

Why Multithreading?

Threads are popular way to improve application through parallelism. For example, in a browser, multiple tabs can be different threads.

Threads operate faster than processes due to following reasons:

- 1) Thread creation is much faster.
- 2) Context switching between threads is much faster.
- 3) Threads can be terminated easily
- 4) Communication between threads is faster.



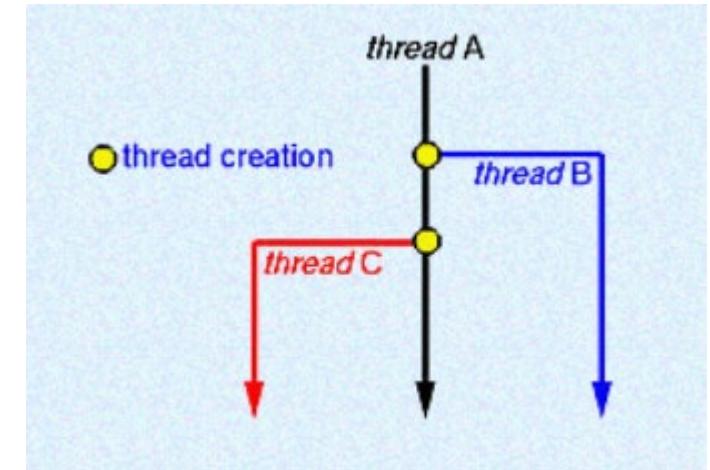
Can we write multithreading programs in C?

Unlike Java, multithreading is not supported by the C language standard. [POSIX Threads \(or Pthreads\)](#) is a POSIX standard for threads. Implementation of pthread is available with gcc compiler.

- To the software developer, the concept of a "procedure" that runs independently from its main program may best describe a thread.
- Pthread: POSIX Thread, a standard-based thread API for C.
- other options: openMP, std::thread
- When compiling Pthread in gcc/g++, should add option “-lpthread”.
 - Compile: `gcc test.c -lpthread` or `g++ test.cpp -lpthread`
 - Execution: `./a.out`

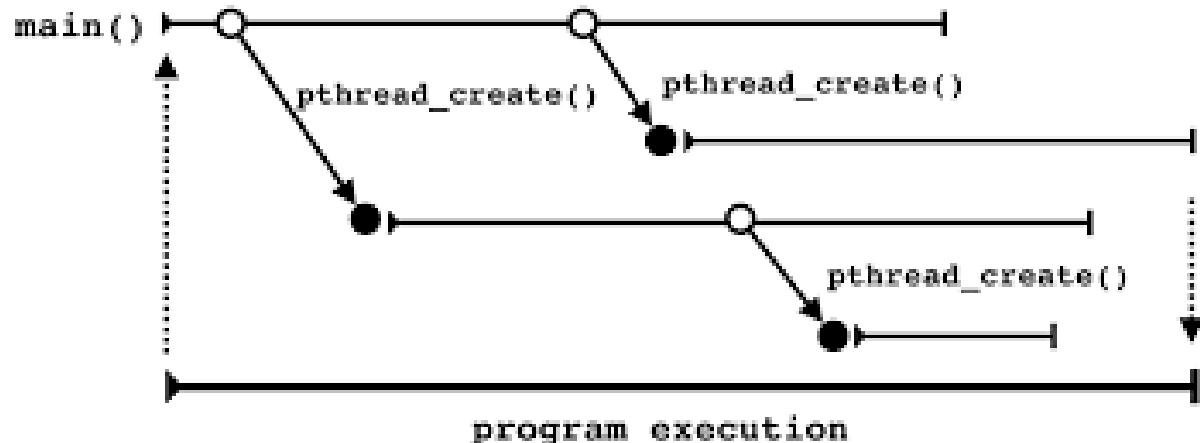
Pthread creation

- **pthread_create:**
 - `int pthread_create(pthread_t *thread,
const pthread_attr_t *attr,
void *(*start_routine) (void *),
void *arg);`
- thread: pointer to an unsigned integer value that returns the thread id of the thread created.
- The attr parameter is used to set thread attributes. You can specify a thread attributes object like scheduling policy, detached state, etc. **Set NULL by default.**
- The start_routine is the C **function** that the thread will execute once it is created.
- arg: pointer to void that contains the **arguments** to the function



Pthread creation

- Return value
 - On success, `pthread_create()` returns 0;
 - On error, it returns an error number, and the contents of `*thread` are undefined.
- Pthread is declared with type:
 - `pthread_t` (defined in “`sys/types.h`”)



Pthread creation

- The execution order of created threads depends on the task scheduling of the operating system
- It might seem like out of order

The image shows a terminal window titled "Terminal" running on a VM. The command `gcc Pthread_Creation.c -lpthread` is run to compile the source code, and then `./a.out` is run to execute it. The output shows the main thread creating five threads (0-4) and then printing "Hello world! thread <tid>\n" for each thread in an interleaved manner.

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 #define NUM_THREAD 5
7
8 void *print_hello(void *threadid){
9     long tid;
10    tid = (long)threadid;
11
12    printf("Hello world! thread %ld\n", tid);
13    pthread_exit(NULL);
14 }
15
16 int main(){
17
18    pthread_t threads[NUM_THREAD];
19    int rc;
20    long i;
21
22    for(i =0; i<NUM_THREAD; i++){
23        printf("In main: create thread %ld\n", i);
24        rc = pthread_create(&threads[i], NULL, print_hello, (void*)i);
25        if(rc){
26            printf("ERROR: return code from pthread_create() is %d", rc);
27            exit(1);
28        }
29    }
30
31
32    pthread_exit(NULL);
33    return 0;
34 }
```

```
[10/09/18]seed@VM:~/.../Pthread_Creation$ gcc Pthread_Creation.c -lpthread
[10/09/18]seed@VM:~/.../Pthread_Creation$ ./a.out
In main: create thread 0
In main: create thread 1
In main: create thread 2
In main: create thread 3
In main: create thread 4
Hello world! thread 4
Hello world! thread 3
Hello world! thread 2
Hello world! thread 1
Hello world! thread 0
[10/09/18]seed@VM:~/.../Pthread_Creation$
```

Pthread termination

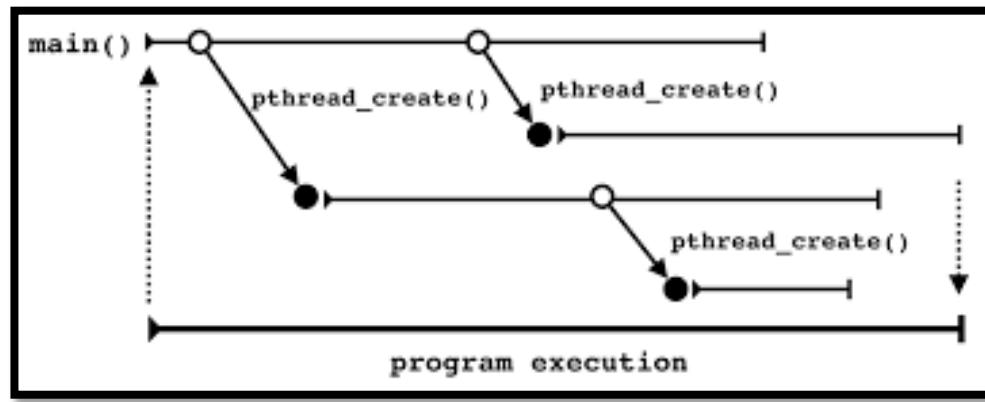
- **pthread_exit:**

- `void pthread_exit(void *retval);`

Parameters: This method accepts a mandatory parameter **retval** which is the pointer to an integer that stores the return status of the thread terminated.

This function is used to **explicitly exit a thread**. Typically, the `pthread_exit()` function is called after a thread has completed its work.

- If `main()` finishes before the **threads it has created, and exits with `pthread_exit()`**, the other threads will **continue** to execute. Otherwise, they will **be automatically terminated** when `main()` finishes.
- Recommendation: Use `pthread_exit()` to exit from all threads...especially `main()`.



Pthread termination – main thread without `pthread_exit()`

Other threads would be automatically terminated when `main()` finishes. Therefore, they might not have finished their work.

The image shows a terminal window titled "Terminal" running on a VM. The command `gcc Pthread_Termination.c -lpthread` is run to compile the program, and then `./a.out` is run to execute it. The output shows that the main thread exits immediately after creating the thread, while the thread itself continues to run and prints its message.

```
Open ▾ +  
1 #include <pthread.h>  
2 #include <stdio.h>  
3 #include <stdlib.h>  
4 #include <unistd.h>  
5  
6  
7 void *print_hello(void *threadid){  
8  
9     sleep(2);  
10    printf("Hello world!\n");  
11    pthread_exit(NULL);  
12 }  
13  
14 int main(){  
15     pthread_t thread;  
16     int rc;  
17     void* i;  
18  
19     printf("In main: create thread\n");  
20     rc = pthread_create(&thread, NULL, print_hello, i);  
21  
22     if(rc){  
23         printf("ERROR: return code from pthread_create() is %d", rc);  
24         exit(1);  
25     }  
26  
27     printf("Main thread exits!\n");  
28     //pthread_exit(NULL);  
29  
30     return 0;  
31 }
```

```
[10/09/18]seed@VM:~/.../Pthread_Termination$ gcc Pthread_Termination.c -lpthread  
[10/09/18]seed@VM:~/.../Pthread_Termination$ ./a.out  
In main: create thread  
Main thread exits!  
[10/09/18]seed@VM:~/.../Pthread_Termination$ █
```

Pthread termination – main with `pthread_exit()`.

The other threads will **continue** to execute.

The terminal window shows the following output:

```
[10/08/18]seed@VM:~/.../Pthread_Termination$ gcc Pthread_Termination.c -lpthread
[10/08/18]seed@VM:~/.../Pthread_Termination$ ./a.out
In main: create thread
Main thread exits!
Hello world!
[10/08/18]seed@VM:~/.../Pthread_Termination$
```

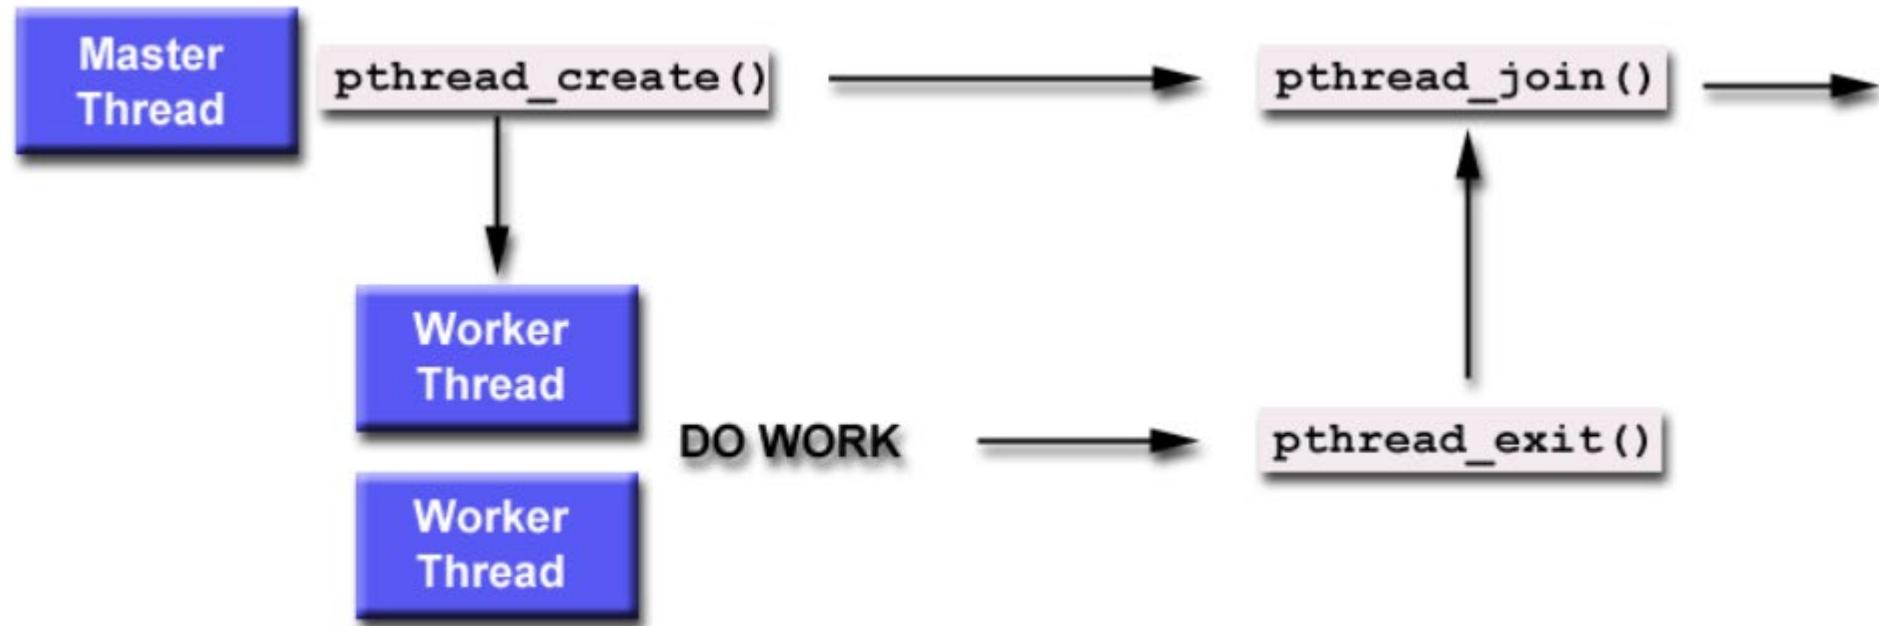
The code in the editor is as follows:

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6
7 void *print_hello(void *threadid){
8
9     sleep(2);
10    printf("Hello world!\n");
11    pthread_exit(NULL);
12 }
13
14 int main(){
15     pthread_t thread;
16     int rc;
17     void* i;
18
19     printf("In main: create thread\n");
20     rc = pthread_create(&thread, NULL, print_hello, i);
21
22     if(rc){
23         printf("ERROR: return code from pthread_create() is %d", rc);
24         exit(1);
25     }
26
27     printf("Main thread exits!\n");
28     pthread_exit(NULL);
29
30     return 0;
31 }
```

Pthread join - synchronization

- `pthread_join`:
 - `int pthread_join(pthread_t thread, void *retval);`
- "Joining" is one way to accomplish ***synchronization*** between threads.
- The `pthread_join()` function **blocks the calling thread** until the **specified thread terminates**.
- The programmer is able to obtain the target thread's termination return status if specified through `pthread_exit()`, in the status parameter.

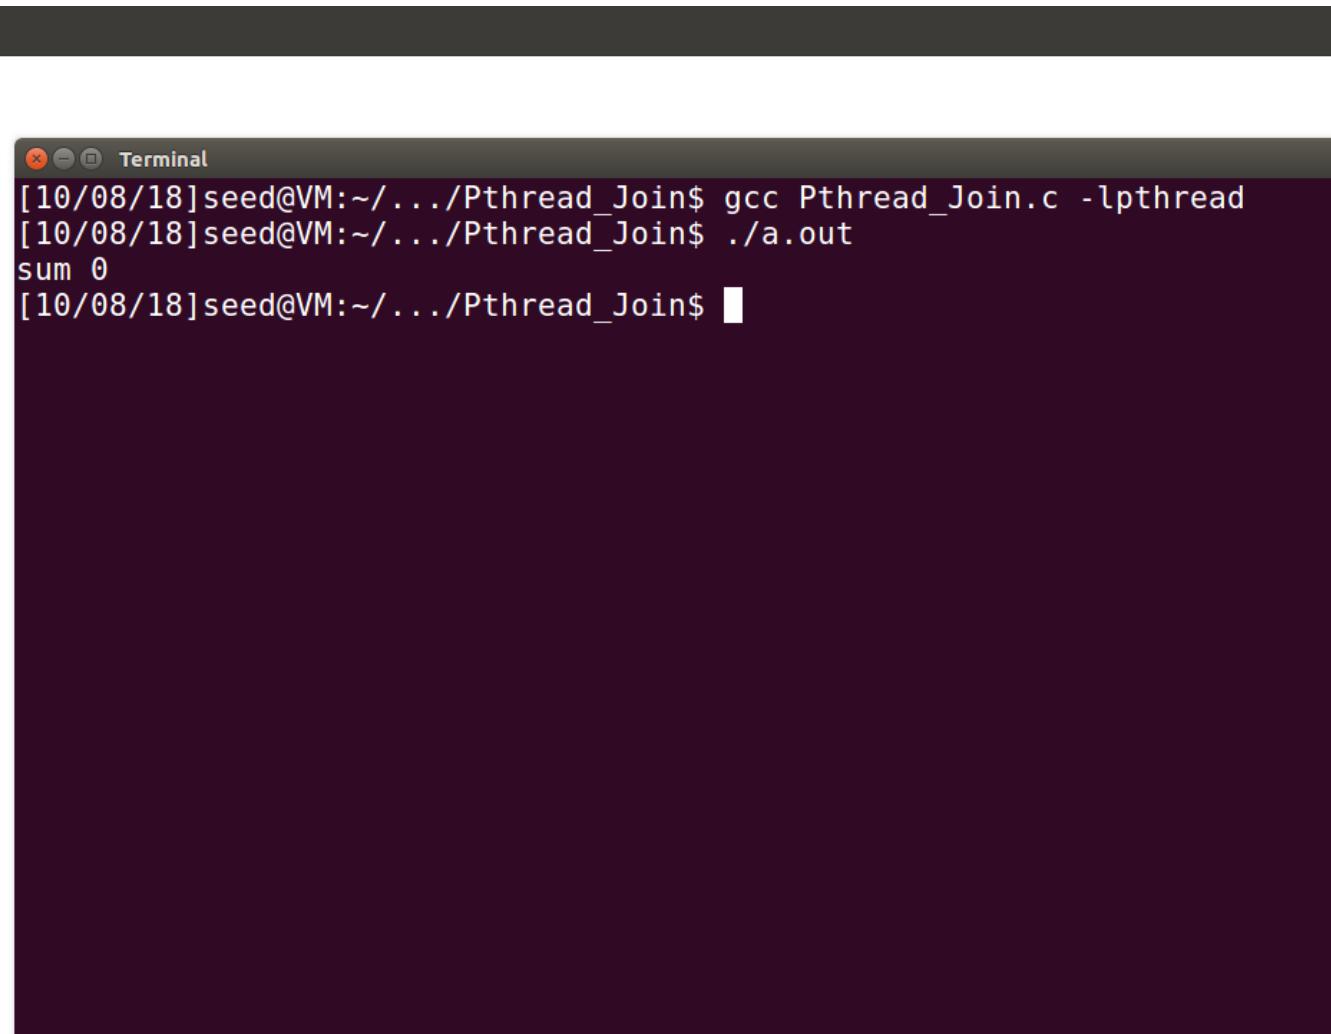
"Joining" is one way to accomplish synchronization between threads.



Pthread join - synchronization

- Return value
 - On success, `pthread_join()` returns **0**;
 - On error, it returns an **error number**.
- It is impossible to join a detached thread.
- When a thread is created, one of its attributes defines whether it is joinable or detached. Detached means it can never be joined. (`PTHREAD_CREATE_DETACHED` or `PTHREAD_CREATE_JOINABLE`)

Pthread join – without calling pthread_join().



```
Open ▾ +  
1 #include<stdlib.h>  
2 #include<stdio.h>  
3 #include<unistd.h>  
4 #include<pthread.h>  
5  
6  
7 int sum;  
8  
9 void * add1(void *cnt)  
10 {  
11     for(int i=0; i < 5; i++)  
12     {  
13         sum += i;  
14     }  
15     pthread_exit(NULL);  
16     return 0;  
17 }  
18 void * add2(void *cnt)  
19 {  
20     for(int i=5; i<10; i++)  
21     {  
22         sum += i;  
23     }  
24     pthread_exit(NULL);  
25     return 0;  
26 }  
27 }  
28  
29 int main(void)  
30 {  
31     pthread_t ptid1, ptid2;  
32     sum=0;  
33  
34     pthread_create(&ptid1, NULL, add1, &sum);  
35     pthread_create(&ptid2, NULL, add2, &sum);  
36  
37     //pthread_join(ptid1,NULL);  
38     //pthread_join(ptid2,NULL);  
39  
40     printf("sum %d\n", sum);  
41     pthread_exit(NULL);  
42  
43  
44     return 0;  
45 }
```

Pthread join – synchronization with calling pthread_join().

The terminal window shows the following output:

```
[10/08/18]seed@VM:~/.../Pthread_Join$ gcc Pthread_Join.c -lpthread
[10/08/18]seed@VM:~/.../Pthread_Join$ ./a.out
sum 45
[10/08/18]seed@VM:~/.../Pthread_Join$
```

The code in the editor is a C program that demonstrates thread synchronization using `pthread_join`. It includes headers for stdlib.h, stdio.h, unistd.h, and pthread.h. The program defines two functions, `add1` and `add2`, which calculate the sum of integer ranges. The main function creates two threads, `ptid1` and `ptid2`, using `pthread_create` and then waits for them to finish using `pthread_join`. The final sum is printed to the console.

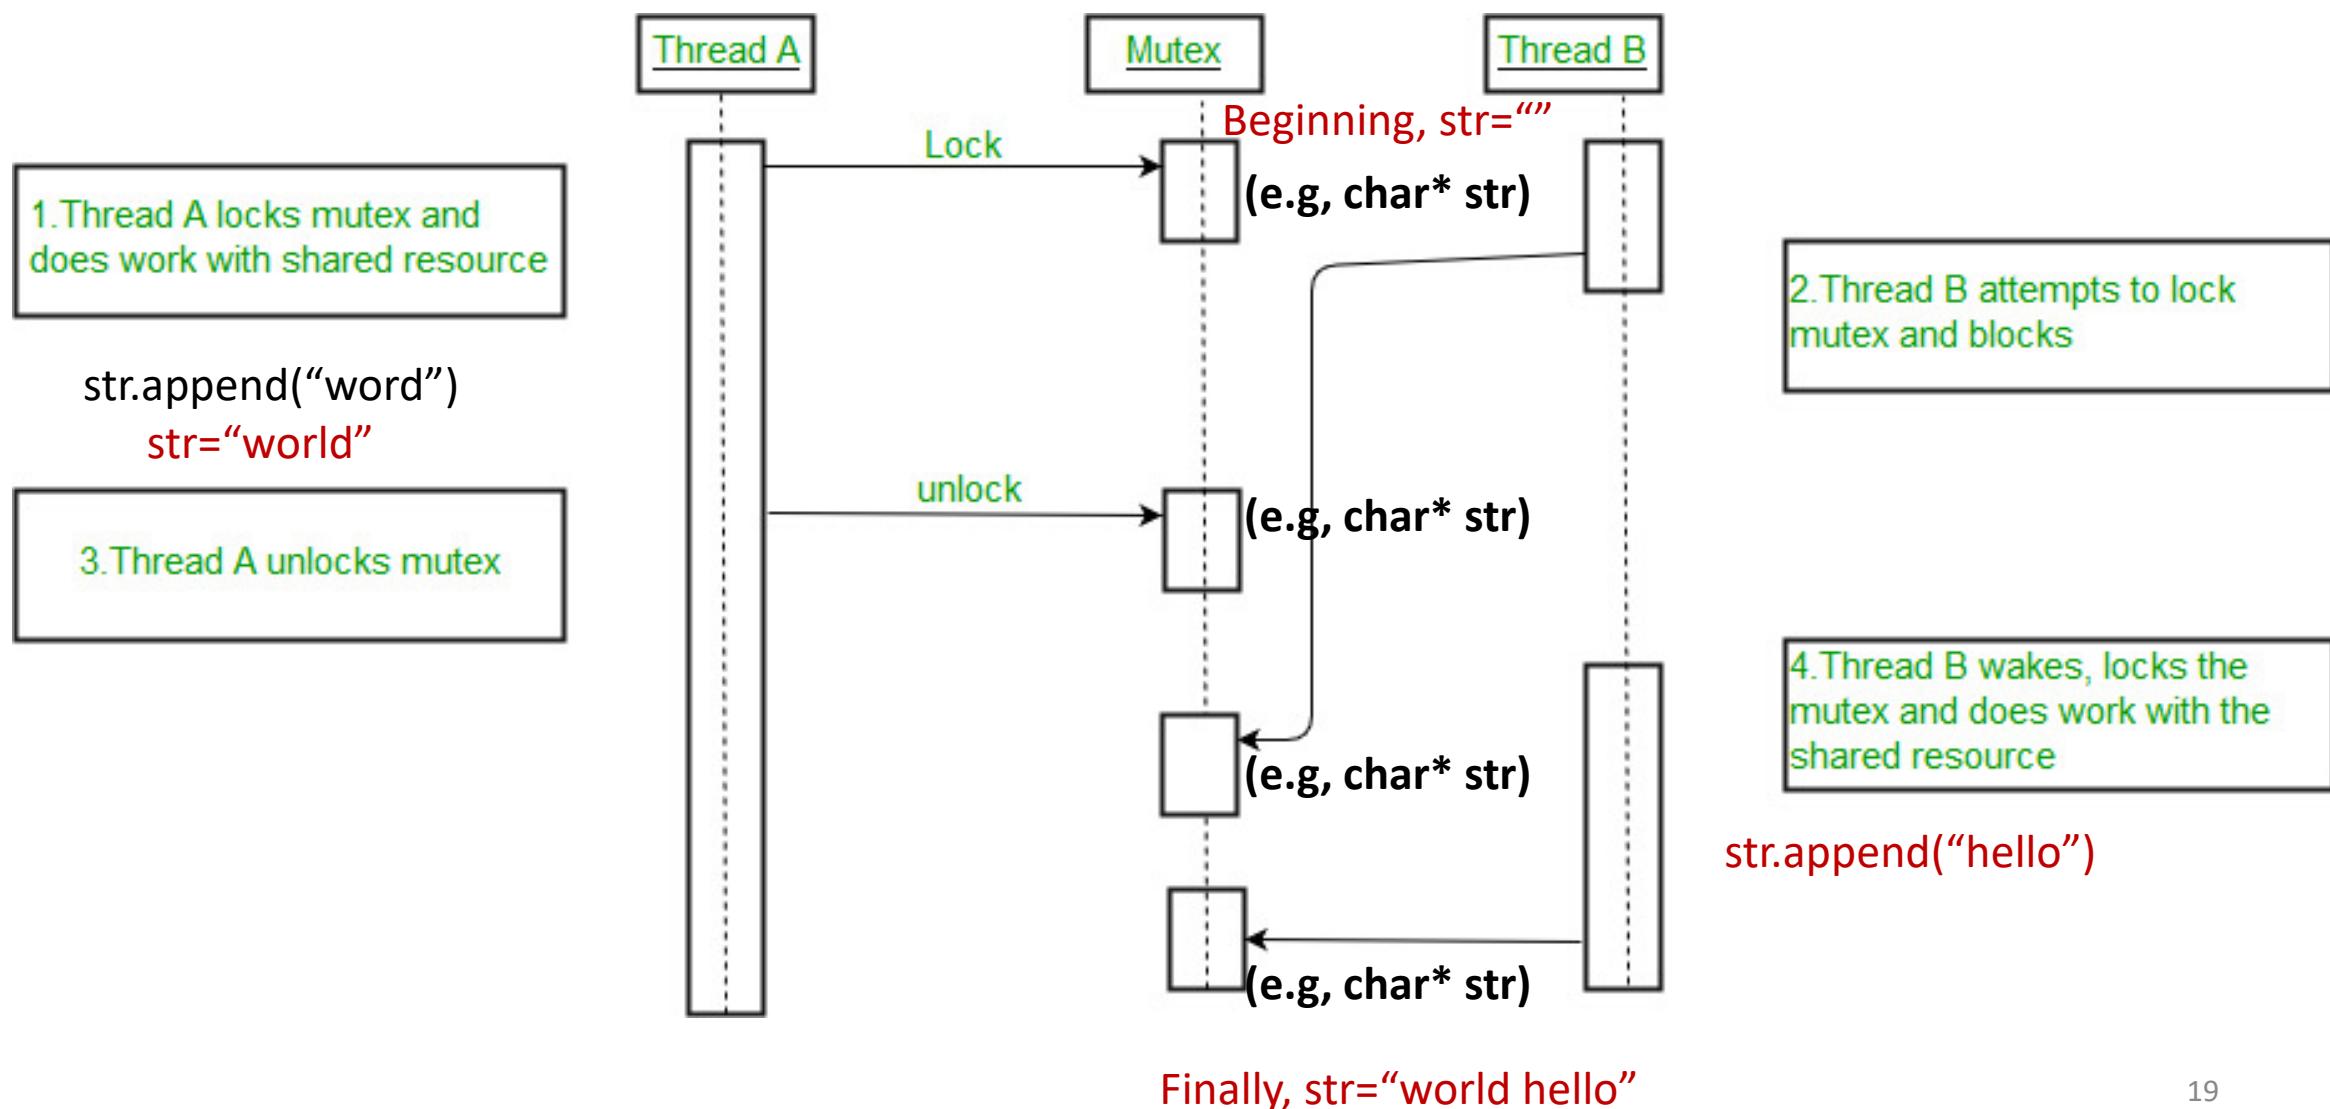
```
1 #include<stdlib.h>
2 #include<stdio.h>
3 #include<unistd.h>
4 #include<pthread.h>
5
6
7 int sum;
8
9 void * add1(void *cnt)
10 {
11     for(int i=0; i < 5; i++)
12     {
13         sum += i;
14     }
15     pthread_exit(NULL);
16     return 0;
17 }
18 void * add2(void *cnt)
19 {
20
21     for(int i=5; i<10; i++)
22     {
23         sum += i;
24     }
25     pthread_exit(NULL);
26     return 0;
27 }
28
29 int main(void)
30 {
31     pthread_t ptid1, ptid2;
32     sum=0;
33
34     pthread_create(&ptid1, NULL, add1, &sum);
35     pthread_create(&ptid2, NULL, add2, &sum);
36
37     pthread_join(ptid1,NULL);
38     pthread_join(ptid2,NULL);
39
40     printf("sum %d\n", sum);
41     pthread_exit(NULL);
42
43     return 0;
44 }
```

The annotation points to the line `pthread_join(ptid2,NULL);` in the code, with the text: "The main thread would stop and wait for ptid1 and ptid2 to finish at line 38."

Pthread mutex – flag for privacy/security

- Mutex is an abbreviation for "**mutual exclusion**". Mutex variables are one of the primary means of implementing **thread synchronization** and for protecting shared data when multiple writes occur.
- A **mutex variable** acts like a "lock" protecting access to a shared data resource.

- A mutex variable acts like a "lock" protecting access to a shared data resource.



Mutex

- A Mutex is a lock that we set before using a shared resource and release after using it.
- When the lock is set, no other thread can access the locked region of code.
- So we see that even if thread B is scheduled while thread A was not done accessing the shared resource. **The code is locked by thread A using mutexes then thread B cannot even access that region of code.**
- So this ensures synchronized access of shared resources in the code.

Pthread mutex - flag

- Mutex should be declared with type:
 - `pthread_mutex_t` (defined in “`sys/types.h`”)
- Mutex should be initialized before it is used:
 - `int pthread_mutex_init(pthread_mutex_t *mutex,
const pthread_mutexattr_t *attr);`
 - It initializes the `mutex` referenced by `mutex` with attributes specified by `attr`.
 - If `attr` is `NULL`, the default mutex attributes are used; the effect is the same as passing the address of a default mutex attributes object.
 - Upon successful initialisation, the state of the mutex becomes initialised and unlocked.
- Mutex should be free if it is no longer used:
 - `int pthread_mutex_destroy(pthread_mutex_t *);`

Pthread mutex - flag

- Pthread mutex lock routines:
 - `int pthread_mutex_lock(pthread_mutex_t *mutex);`
 - `int pthread_mutex_trylock(pthread_mutex_t *mutex);`
 - `int pthread_mutex_unlock(pthread_mutex_t *mutex);`

Pthread mutex - flag

- The **pthread_mutex_lock()** routine is used by a thread **to acquire a lock** on the specified mutex variable. **If the mutex is already locked by another thread, this call will block the calling thread until the mutex is unlocked.**
- **pthread_mutex_trylock()** will attempt to lock a mutex. However, if the mutex is already locked, the routine will return immediately with a "busy" error code. This routine may be useful in preventing deadlock conditions, as in a priority-inversion situation.
- **pthread_mutex_unlock()** will unlock a mutex if called by the owning thread. Calling this routine is required after a thread has completed its use of protected data if other threads are to acquire the mutex for their work with the protected data. An **error** will be returned if:
 - **If the mutex was already unlocked**
 - **If the mutex is owned by another thread**

Suppose we have two jobs to update a shared variable called counter.

```
11 ↵ pthread_t tid[2];
12 ↴ int counter;
13
14 void* trythis(void* arg)
15 {
16     unsigned long i = 0;
17     counter += 1;
18     printf( format: "\n Job %d has started\n", counter);
19
20     for (i = 0; i < (0xFFFFFFFF); i++)
21         ;
22     printf( format: "\n Job %d has finished\n", counter);
23     return NULL;
24 }
25 ➤ int main(void)
26 {
27     int i = 0;
28     int error;
29     counter=0;
30     while (i < 2) {
31         error = pthread_create(&(tid[i]), attr: NULL, &trythis, arg: NULL);
32         if (error != 0)
33             printf( format: "\nThread can't be created : [%s]", strerror(error));
34         i++;
35     }
36     pthread_join( th: tid[0], thread_return: NULL);
37     pthread_join( th: tid[1], thread_return: NULL);
```

Line 20 consume more time (e.g. 3 seconds)

Execution Order: A B C D E F

Thread 1 (execute first)

A Line 17: counter=1

B Line 20: much time

E Line 22: J“Job 1”

Thread 2

C Line 17: counter=2

D Line 20: much time

F Line 22: J“Job 2”

It's not what we want!

Run: pthread_without_mutex

```
Job 1 has started
Job 2 has started
Job 2 has finished
Job 2 has finished
Process finished with exit code 0
```

```

13 void* trythis(void* arg)
14 {
15     pthread_mutex_lock(&lock);
16     unsigned long i = 0;
17     counter += 1;
18     printf( format: "\n Job %d has started\n", counter);
19
20     for (i = 0; i < (0xFFFFFFFF); i++)
21         ;
22     printf( format: "\n Job %d has finished\n", counter);
23
24     pthread_mutex_unlock(&lock);
25
26     return NULL;
27 }

33 int main(void)
34 {
35     int i = 0;
36     int error;
37     counter=0;
38     if (pthread_mutex_init(&lock, mutexattr: NULL) != 0) {
39         printf( format: "\n mutex init has failed\n");
40         return 1;
41     }
42     while (i < 2) {
43         error = pthread_create(&(tid[i]),
44                               attr: NULL,
45                               &trythis, arg: NULL);
46         if (error != 0)
47             printf( format: "\nThread can't be created :[%s]",
48                     strerror(error));
49         i++;
50     }
51
52     pthread_join( th: tid[0], thread_return: NULL);
53     pthread_join( th: tid[1], thread_return: NULL);
54     pthread_mutex_destroy(&lock);

```

Line 20 consume more time (e.g. 3seconds)

Execution Order: A B C D E F (impossible)

Execution Order: A B E C D F or C D F A B E

Thread 1 (execute first)

Thread 2

A Line 17: counter=1

C Line 17: counter=2

B Line 20: much time

D Line 20: much time

E Line 22: J“Job 1”

F Line 22: J“Job 2”

It's correct!

Job 1 has started

Job 1 has finished

Job 2 has started

Job 2 has finished

Process finished with exit code 0

Pthread condition - signals

- **Condition variables** provide yet another way for threads to synchronize.
- While **mutexes** implement synchronization by controlling thread access to data, condition variables allow threads **to synchronize based upon the actual value of data**.
- A condition variable is always **used in conjunction with a mutex lock**.

References

- <https://www.baeldung.com/cs/async-vs-multi-threading>
- <https://www.javatpoint.com/process-vs-thread>
- <https://www.geeksforgeeks.org/multithreading-c-2/>
- <https://www.geeksforgeeks.org/condition-wait-signal-multi-threading/>
- http://www.cs.unibo.it/~ghini/didattica/sistop/pthreads_tutorial/POSIX_Threads_Programming.htm
- <https://www.geeksforgeeks.org/mutex-lock-for-linux-thread-synchronization/>

Pthread condition - signals

- Condition variables must be declared with type: `pthread_cond_t`
 - `pthread_cond_t` (defined in “`sys/types.h`”)
- Condition variables must be initialized before it is used:
 - `int pthread_cond_init(pthread_cond_t *, const pthread_condattr_t *);`
- Condition variables should be freed if it is no longer used:
 - `int pthread_cond_destroy(pthread_cond_t *);`

Pthread condition - signals

- Pthread condition routines:

- `int pthread_cond_wait(pthread_cond_t *, pthread_mutex_t *);`
- `int pthread_cond_signal(pthread_cond_t *);`
- `int pthread_cond_broadcast(pthread_cond_t *);`

Pthread condition - signals

- **pthread_cond_wait()** blocks the calling thread until **the specified condition** is signalled. This routine should be called **while mutex is locked**, and it will **automatically release the mutex while it waits**.
- **The pthread_cond_signal()** routine is used **to signal (or wake up) another thread which is waiting** on the condition variable. It should be called after mutex is locked, and must unlock mutex in order for pthread_cond_wait() routine to complete.
- **The pthread_cond_broadcast()** routine should be used instead of pthread_cond_signal() if more than one thread is in a blocking wait state.

Pthread condition - signals

```
Open ▾
```

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <unistd.h>
4
5 #define NUM_THREADS 3
6 #define TCOUNT      10
7 #define COUNT_LIMIT 10
8
9 int count = 0;
10 int thread_ids[3] = {0,1,2};
11 pthread_mutex_t count_mutex;
12 pthread_cond_t  count_threshold_cv;
13
14 void *inc_count(void *idp)
15 {
16     int i = 0;
17     int taskid = 0;
18     int *my_id = (int*)idp;
19
20     for (i=0; i<TCOUNT; i++) {
21         pthread_mutex_lock(&count_mutex);
22         taskid = count;
23         count++;
24
25         if (count == COUNT_LIMIT){
26             pthread_cond_signal(&count_threshold_cv);
27         }
28
29         printf("inc_count(): thread %d, count = %d, unlocking mutex\n", *my_id, count);
30         pthread_mutex_unlock(&count_mutex);
31         sleep(1);
32     }
33
34     printf("inc_count(): thread %d, Threshold reached.\n", *my_id);
35
36     pthread_exit(NULL);
37 }
```

```
39 void *watch_count(void *idp)
40 {
41     int *my_id = (int*)idp;
42     printf("Starting watch_count(): thread %d\n", *my_id);
43
44     pthread_mutex_lock(&count_mutex);
45
46     while(count<COUNT_LIMIT) {
47         pthread_cond_wait(&count_threshold_cv, &count_mutex);
48         printf("watch_count(): thread %d Condition signal received.\n", *my_id);
49     }
50
51     count += 100;
52     pthread_mutex_unlock(&count_mutex);
53     pthread_exit(NULL);
54 }
55
56 int main (int argc, char *argv[])
57 {
58     int i, rc;
59     pthread_t threads[3];
60     pthread_attr_t attr;
61
62     /* Initialize mutex and condition variable objects */
63     pthread_mutex_init(&count_mutex, NULL);
64     pthread_cond_init (&count_threshold_cv, NULL);
65
66     /* For portability, explicitly create threads in a joinable state */
67     pthread_attr_init(&attr);
68     pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
69     pthread_create(&threads[0], &attr, inc_count, (void *)&thread_ids[0]);
70     pthread_create(&threads[1], &attr, inc_count, (void *)&thread_ids[1]);
71     pthread_create(&threads[2], &attr, watch_count, (void *)&thread_ids[2]);
72
73     /* Wait for all threads to complete */
74     for (i=0; i<NUM_THREADS; i++) {
75         pthread_join(threads[i], NULL);
76     }
77     printf ("Main(): Waited on %d threads. Done.\n", NUM_THREADS);
78
79     /* Clean up and exit */
80     pthread_attr_destroy(&attr);
81     pthread_mutex_destroy(&count_mutex);
82     pthread_cond_destroy(&count_threshold_cv);
83     pthread_exit(NULL);
84
85     return 0;
86 }
```

Release the mutex at line 47

Pthread condition - signals

```
[10/09/18]seed@VM:~/.../Pthread_Cond$ gcc Pthread_Cond.c -lpthread
[10/09/18]seed@VM:~/.../Pthread_Cond$ ./a.out
Starting watch_count(): thread 2
inc_count(): thread 1, count = 1, unlocking mutex
inc_count(): thread 0, count = 2, unlocking mutex
inc_count(): thread 1, count = 3, unlocking mutex
inc_count(): thread 0, count = 4, unlocking mutex
inc_count(): thread 1, count = 5, unlocking mutex
inc_count(): thread 0, count = 6, unlocking mutex
inc_count(): thread 1, count = 7, unlocking mutex
inc_count(): thread 0, count = 8, unlocking mutex
inc_count(): thread 1, count = 9, unlocking mutex
inc_count(): thread 0, count = 10, unlocking mutex
watch_count(): thread 2 Condition signal received.
inc_count(): thread 1, count = 111, unlocking mutex
inc_count(): thread 0, count = 112, unlocking mutex
inc_count(): thread 1, count = 113, unlocking mutex
inc_count(): thread 0, count = 114, unlocking mutex
inc_count(): thread 1, count = 115, unlocking mutex
inc_count(): thread 0, count = 116, unlocking mutex
inc_count(): thread 1, count = 117, unlocking mutex
inc_count(): thread 0, count = 118, unlocking mutex
inc_count(): thread 1, count = 119, unlocking mutex
inc_count(): thread 0, count = 120, unlocking mutex
inc_count(): thread 1, Threshold reached.
inc_count(): thread 0, Threshold reached.
Main(): Waited on 3 threads. Done.
[10/09/18]seed@VM:~/.../Pthread_Cond$ █
```

Thank you